

# PDSS CW2

B174502

## Real-Time Distributed Fraud Detection System

### Input

```
case class Transaction(transactionID: String, fromID: String, toID: String, amount: Float,
                        timestamp: Long)
RDD[Transaction]
```

Where each entry represents a transaction that is potentially fraudulent. The `timestamp` is a `Long` and is represented as the number of seconds since the Unix epoch

### Output

```
RDD[Transaction]
```

Where each entry is a subset of the input containing the fraudulent transactions in question with a confidence score

**Solution** We can decide whether a transaction is fraudulent based on the transaction amount and the frequency of transactions

1. Define entry function

```
def detect_fraudulent_transactions(transactionsRDD: RDD[Transaction]): RDD[Transaction]
```

2. Reduce the input RDD to calculate the sum and squared sums of transaction amounts

```
val transaction_amount_stats = transactionsRDD
    .map(x => (x.amount, x.amount*x.amount, 1))
    .reduce((x,y) => (x._1 + y._1, x._2 + y._2, x._3 + y._3))
```

3. Calculate mean and standard deviation of transaction amounts

```
/// Obvious calculation
```

4. From each transaction, do a self-join (cartesian join) to get all combinations of transactions and filter to get a strict ordering of transactions ( $T1 < T2$ ) and map to get the time between transactions

```
val transaction_time = transactionsRDD.map(x => (x.transactionId, x.timestamp))
val transaction_time_diff = transaction_time.cartesian(transaction_time)
    .filter(x => x._1._2 < x._2._2)
    .map(x => (x._2._1, x._2._2 - x._1._2))
```

5. Count the number of transactions within a minute window

```
val transaction_time_diff_in_min = transaction_time_diff
.map{case x => if (x._2 < 60) {(x._1,1)} else {(x._1,0)}}
.reduceByKey((x,y) => x+y)
```

7. Calculate the mean and standard deviation of the count of transactions within 1-minute

```
val transaction_time_stats_mean = transaction_time_diff_in_min
.map(x => x._2)
.reduce((x,y) => x + y) / transaction_time_diff_in_min.count()
```

```
val transaction_time_stats_std = math.sqrt(
transaction_time_diff_in_min.map(x => (x._2 - transaction_time_stats_mean)*(x._2 - transaction_time_stats_mean))
.reduce((x,y) => x + y) / transaction_time_diff_in_min.count()
)
```

9. Determine out-of-distribution transaction frequency and transaction amounts

```
val fraud_time = transactionsRDD
.map(x => (x.transactionId, x)).join(transaction_time_diff_in_min)
.filter(x => x._2._2 > transaction_time_stats_mean + 1*transaction_time_stats_std)
.map(x => x._2._1)
val fraud_amount = transactionsRDD
.filter(x => x.amount > transaction_amount_mean + 1*transaction_amount_std)
val fraud = fraud_time.union(fraud_amount).distinct()
```

## Optimisations

- Persist RDDs by caching intermediate RDDs that are used across calculations
- Partition `transactionsRDD` with `transactionID` prior to first operation and then the final joins to get the fraudulent transactions
- Use an alternative method than `cartesian` to avoid a large join, perhaps a custom join that only joins the transactions close together in timestamp
  - We can only use range partitioning to partition this dataset according to the timestamp

## Task 2

To cross-reference suspicious activities we can consider both location history for transactions. This can be done by considering a directed graph to find and suspicious transactions

1. Initialize using the entire transaction history
  - (a) Initialise a node for each user
  - (b) Initialise edges as the transaction with their edge weights initialised with:
    - i. Transaction Frequency
    - ii. Transaction Amounts
  - (c) Initialise a clustered location graph for each user
    - i. Cluster user locations based on -means
    - ii. Each edge is a transaction from two different locations
2. User and Transaction Pattern Detection
  - (a) Cycle Detection
    - i. Find cycles within the graph where they are completed within a short-time frame
  - (b) Flow Analysis
    3. Money flow into specific accounts that are unusual
    4. Splitting and merging of money across nodes
  - (c) Node Metrics
    - i. Look at if transaction connects to known suspicious users

### 3. Location Pattern Detection

- (a) Whether the newly created transaction has a high confidence score for being added to a location cluster

## Distributed Log Processing for Anomaly Detection

### Input

```
case class LogEntry(timestamp: Long, eventType: String, processId: String, responseTime: Long)
RDD[LogEntry]
```

The input is an RDD where each entry is a log where they contain the timings of the log (in seconds since Unix epoch), the type of the log and the process that generated the log.

### Output

We can output all `LogEntry` which we consider to be unusual or suspicious. Since we will be adopting a window-based aggregation approach. All logs within a specific window will be outputted

```
RDD[(LogEntry)]
```

### Solution

We should detect whether there are anomalies in logs based on the following:

- Sudden spikes in error rates
- Request time degradation across the same process

1. Define the entry function

```
def detect_unusual_log_patterns(logs: RDD[LogEntry]): RDD[LogEntry]
```

2. Map the RDD to transform timestamps to its corresponding window

```
val windowed_logs = logs.map(x => (x.timestamp / WINDOW_SIZE, x))
```

3. Detect error rate spike anomalies

4. Aggregate error count in each window and so, calculate the error rate in each window

```
val error_counts = windowed_logs
  .map(x => (x._1, (if (x._2.eventType == "ERROR") 1 else 0, 1)))
  .reduceByKey((x,y) => (x._1 + y._1, x._2 + y._2)).mapValues(x => (x._1.toDouble / x._2.toDouble))
```

5. Calculate the mean and standard deviation of the error rate

```
val error_count_mean = error_counts.map(x => x._2).mean()
```

```
val error_count_std = error_counts.map(x => x._2).stdev()
```

6. Detect the out-of-distribution error rate logs

```
val unusual_error_logs = error_counts
  .filter(x => x._2 > error_count_mean + 2 * error_count_std)
  .join(windowed_logs).map(x => x._2._2)
```

7. Detect response time anomalies

8. Reduce by the process key and calculate the mean and standard deviation of response times for this process

```

val process_logs = logs
  .map(x => (x.processId, x)).groupByKey()
  .mapValues(x => {
    val mean = x.map(_.responseTime).reduce(_ + _) / x.size
    val std = math.sqrt(
      x.map(y => math.pow(y.responseTime - mean, 2)).reduce(_ + _) / x.size
    )
    (mean, std)
  })

```

9. Detect the out-of-distribution response-time logs for either high or low response-times

```

val unusual_process_logs = process_logs.filter(x => x._2._2 > x._2._1 * 0.5)
  .join(logs.map(x => (x.processId, x))).map(x => x._2._2)
val unusual_logs = unusual_error_logs.union(unusual_process_logs).distinct()

```

### Optimisations

- Persist intermediate RDDs by saving it in memory using `.persist()`
- Range partition by timestamps so that we can efficiently fetch windowed data in order while not being explicitly hashable

## Task 2

1. Prepare log entries to be used as a timeseries
  - (a) Group into time-based windows
  - (b) Aggregate error counts and types within the windows
  - (c) Baseline statistics for each window
2. Detect spikes in error rates
  - (a) Calculate rolling statistics from window to window for error rates
  - (b) Use statistical methods to identify strong deviations in error rate
  - (c) Flag windows that are sufficiently deviated from typical behaviour